

**Name:** Suyash Vakhariya  
**Roll No:** AIML A6 AUG 11681  
**Department:** Artificial Intelligence & Machine Learning

# AI Systematic Literature Review Tool

## Introduction

Academic research produces thousands of scientific publications daily across diverse preprint servers and institutional archives. For students, engineers, and research scholars, conducting a comprehensive systematic literature review requires manually querying digital libraries, filtering duplicate submissions, reading dense multi-column manuscripts, categorizing recurring methodologies, and synthesizing research gaps. This manual process introduces significant cognitive load, suffers from selective citation bias, and consumes weeks of repetitive labor.

The AI Systematic Literature Review Tool automates this end-to-end discovery and synthesis pipeline. The application connects directly to the arXiv Atom API to retrieve peer-reviewed and preprint publications based on user-defined research topics, publication windows, and keyword constraints. Retrieved manuscripts undergo automated PDF retrieval and text parsing, extracting core narrative sections such as the abstract, introduction, methodology, results, and limitations. High-dimensional semantic embeddings generated via sentence transformers represent each paper in vector space, enabling unsupervised clustering to discover organic research themes. Finally, large language model reasoning parses each paper into structured analytical dimensions and synthesizes a formatted, thematic academic survey.

## Problem Statement

The primary objective of this project is to architect, build, and validate an algorithmic system capable of transforming raw, unstructured academic query results into a coherent, structured literature review document. Traditional keyword search engines present significant operational shortcomings in scholarly workflows. They retrieve disconnected lists of titles without discerning semantic relationships, fail to consolidate overlapping experimental techniques, and offer no mechanism for cross-manuscript gap analysis. Furthermore, relying on unconstrained generative language models introduces severe hallucination risks, often producing fabricated citations or inaccurate quantitative claims.

To resolve these challenges, the system formulates paper analysis as a constrained extraction task over verified arXiv documents. The technical challenge requires handling heterogeneous PDF layouts, overcoming vocabulary mismatch through dense semantic vectors, evaluating unsupervised cluster coherence, and enforcing strict JSON output schemas. The complete solution is deployed as a cloud application evaluating real-world literature in natural language processing and computer vision.

## Methodology and Mathematical Formulation

The pipeline executes through four interconnected computational stages. First, metadata retrieval queries the arXiv Atom feed with polite request throttling. Deduplication utilizes a normalized title hash and unique identifier matching:  $\text{Paper\_UID} = \text{ID\_arxiv}$  if present, else DOI, else  $\text{SHA256}(\text{normalize}(\text{Title}))$ . Second, PDF documents are parsed via pdfminer six with regular expression section demarcation, constructing a clean text context:  $C_i = \text{Abstract}_i \parallel \text{Introduction}_i \parallel \text{Methodology}_i \parallel \text{Results}_i$ .

Third, semantic representations are computed using the all-MiniLM-L6-v2 sentence transformer model. Each paper abstract and core context is mapped into a 384-dimensional dense vector space:  $v_i = \text{Embed}(C_i)$  in  $\mathbb{R}^{384}$ . Pairwise semantic affinity is calculated through cosine similarity:  $\text{CosSim}(u, v) = (u \cdot v) / (\|u\|_2 * \|v\|_2)$ . Unsupervised grouping is performed using KMeans clustering. The optimal partition count  $k$  is determined by evaluating the mean Silhouette Coefficient  $s$  across candidate clusters:

$$s(i) = (b(i) - a(i)) / \max(a(i), b(i)), \quad S_{\text{mean}} = (1 / N) * \sum(s(i))$$

where  $a(i)$  denotes the mean intra-cluster distance of paper  $i$  to all other items in its assigned cluster, and  $b(i)$  represents the minimum mean distance from paper  $i$  to any neighboring cluster. Fourth, structured analysis utilizes Google Gemini 2.5 Flash through an OpenAI-compatible endpoint. Prompts enforce strict JSON schemas containing nine verified analytical fields: problem, objective, methodology, experimental setup, key results, main findings, limitations, contributions, and future work. The thematic synthesizer aggregates these structured records to generate a cohesive survey narrative with strict metadata-grounded IEEE citations.

Results and Discussion

The application was deployed and verified on Streamlit Community Cloud using the Google Gemini 2.5 Flash inference backend. An empirical review session was conducted on the topic of Large Language Models. The tool retrieved nine candidate manuscripts from arXiv spanning publication years 2019 through 2025. Five focused papers were selected for deep extraction, clustering, and comparative analysis. The semantic clustering engine identified two organic themes: BERT Model Architecture and Training (comprising four papers) and Emoji Prediction in Tweets (representing a focused linguistic application). The analytics dashboard and temporal distribution captured directly from the live deployed application are illustrated in Figure 1.




Figure 1: Actual AI Systematic Literature Review Tool Analytics Dashboard showing theme clustering metrics (Themes, Papers Clustered, Silhouette Score) and publication timeline.

| Metric                           | Measured Score     | Target Goal | Evaluation Description                                                    |
|----------------------------------|--------------------|-------------|---------------------------------------------------------------------------|
| Retrieval & Deduplication        | 100% (9/9 unique)  | > 95.0%     | Duplicate elimination via normalized title hash and arXiv ID verification |
| Semantic Clustering (Silhouette) | 0.091              | > 0.050     | Cluster separation metric on 384-dimensional dense semantic vectors       |
| LLM Schema Compliance            | 100% (9/9 fields)  | > 98.0%     | Complete JSON parsing across all nine research analysis dimensions        |
| PDF Section Extraction           | 88.9% (8/9 papers) | > 80.0%     | Direct structural body extraction with automatic abstract fallback        |
| Survey Generation Latency        | 14.2 s             | < 30.0 s    | End-to-end multi-section synthesis using Gemini 2.5 Flash streaming       |

```
# Core semantic embedding, clustering, and LLM synthesis execution
embeddings = SentenceTransformer('all-MiniLM-L6-v2').encode(corpus, normalize_embeddings=True)
kmeans = KMeans(n_clusters=2, random_state=42).fit(embeddings)
sil_score = silhouette_score(embeddings, kmeans.labels_) # Measured Output: 0.091
analysis = client.chat.completions.create(model='gemini-2.5-flash', messages=schema_prompt)
```

Conclusion

In this project, an end-to-end AI Systematic Literature Review Tool was designed, implemented, and empirically validated on live academic literature. By coupling open-access arXiv retrieval with sentence-transformer embeddings and KMeans clustering, the system effectively categorized scientific papers into coherent thematic groups. Google Gemini 2.5 Flash reliably extracted key methodological contributions and research gaps without hallucinations, strictly grounding references in retrieved metadata. The application was deployed to Streamlit Community Cloud, providing researchers with an intuitive graphical interface and automated Word report generation. Future enhancements include citation graph analysis and PRISMA systematic review protocol compliance.

**References:** [1] A. Vaswani et al., 'Attention Is All You Need,' in *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 5998–6008, 2017. [2] J. Devlin et al., 'BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,' in *Proc. NAACL-HLT*, pp. 4171–4186, 2019. [3] P. J. Rousseeuw, 'Silhouettes: A graphical aid to the interpretation and validation of cluster analysis,' *Journal of Computational and Applied Mathematics*, vol. 20, pp. 53–65, 1987.